

Wenrui (Raymond) Wang

Vancouver, CA | raymond.w.wang@hotmail.com | 647-680-3933

| [linkedin.com/in/raymondwwang](https://www.linkedin.com/in/raymondwwang) | github.com/raymondww | <https://raymondww.github.io/>

SUMMARY

Data Scientist with 3+ years of industry experience building production ML systems — LLM pipelines, propensity scoring, applied NLP, and computer vision — at JLL and UBC. Comfortable across the full stack from data ingestion and HPC-distributed model training to business-facing deployment. Currently building a RAG search engine from scratch and actively exploring LLM application development.

TECHNICAL SKILLS

Languages: Python, R, SQL

ML / CV: PyTorch, YOLO (Ultralytics), OpenCV, scikit-learn, Hugging Face, LLM APIs

Data & Cloud: Apache Airflow, Spark, AWS S3, Azure, Alibaba Cloud; PostgreSQL, MySQL, MongoDB, PolarDB

Infrastructure: Docker, Git, SLURM, HPC distributed training

Visualisation & Reporting: Tableau, Power BI, Quarto, MkDocs

EXPERIENCE

MooVision — Computer Vision Capstone — UBC Animal Welfare Program Vancouver, Canada | Apr 2026 – June 2026

- Built MooVision, a computer vision pipeline for automated cross-sucking behaviour detection in dairy calves, fine-tuning YOLO on 1,100 annotated clips across 8 experimental splits (random, time-based, pen-based, period-based); fine-tuned model achieved 0.73 avg temporal IoU vs 0.70 baseline, indicating accurate event localisation in a naturalistic low-prevalence detection setting
- Engineered the full pipeline from 300GB of raw barn footage to automated evaluation reports — OpenCV clipping tool, proximity-based YOLO baseline, Seq-NMS temporal linking, and a Makefile workflow integrating preprocessing, SLURM-parallelised HPC training, inference, and Quarto report generation

Data Scientist II — JLL (Jones Lang LaSalle)

Guangzhou, China | Jul 2023 – Aug 2024

- Architected an end-to-end production ETL pipeline ingesting 9 real estate news outlets via adversarial scraping, orchestrated daily on Alibaba Cloud via Apache Airflow with PolarDB for query serving — replacing a manual weekly research process and delivering city-level news recommendations to external firms in real time
- Built the NLP layer on top of the pipeline applying LLM summarisation, NER-based location and monetary entity extraction, and cosine similarity deduplication to rank and surface relevant news by city; reduced duplicate content surfaced to analysts and enabled real-time delivery at scale

Data Scientist — JLL (Jones Lang LaSalle)

Guangzhou, China | Jan 2023 – Jul 2023

- Designed a multi-signal propensity scoring pipeline fusing corporate signals (job postings, investment flows, policy changes) with scraped POI data to predict client upsizing/downsizing intent, surfacing prioritised leads via a context-aware recommendation system; contributed to a 30% increase in broker cold-call success rate
- Led a GPT-3.5 hackathon project: built a Flask platform using Selenium-driven dynamic rendering and runtime LLM-generated code execution to extract structured data from arbitrary websites, enabling zero-code data collection for non-technical analysts
- Built a street-view imagery scraper with geo-coordinate extraction to map land plot visual context for an Industry Park white paper, delivering proprietary macro analysis of development pace and land usage ahead of public reporting

Business Data Analyst — JLL (Jones Lang LaSalle)

Guangzhou, China | Sep 2021 – Jan 2023

- Co-authored the Vacancy Anxiety Index: engineered a quarter-aware vacancy tracking algorithm across 4 submarkets, 21 buildings, and 500+ units, combining per-unit destocking duration with an absorption rate model to produce a forward-looking leasing risk indicator used directly in institutional client advisory
- Built IQR-based pricing benchmarks integrating third-party market data to flag statistical outliers in landlord asking rents, translating anomaly detection into premium positioning recommendations for the leasing advisory team
- Automated legacy BI reporting in Python and Tableau and built ad-hoc POI scraping pipelines for new market entry, eliminating manual reporting errors and accelerating regional business development

EDUCATION

Master of Data Science University of British Columbia

Aug 2025 – Jun 2026

MicroMasters — Statistics & Data Science MIT Open Learning

May 2023 – Jan 2025

Bachelor of Commerce — Business Technology Management

Sep 2016 – Jun 2021

Toronto Metropolitan University (formerly Ryerson University)